

Martin Mundt, Dr. Iuliia Pliushch, Prof. Dr. Visvanathan Ramesh

Pattern Analysis & Machine Intelligence

Praktikum: MLPR-WS19

Week 12: Introduction into Reinforcement Learning



Reinforcement Learning: Applications


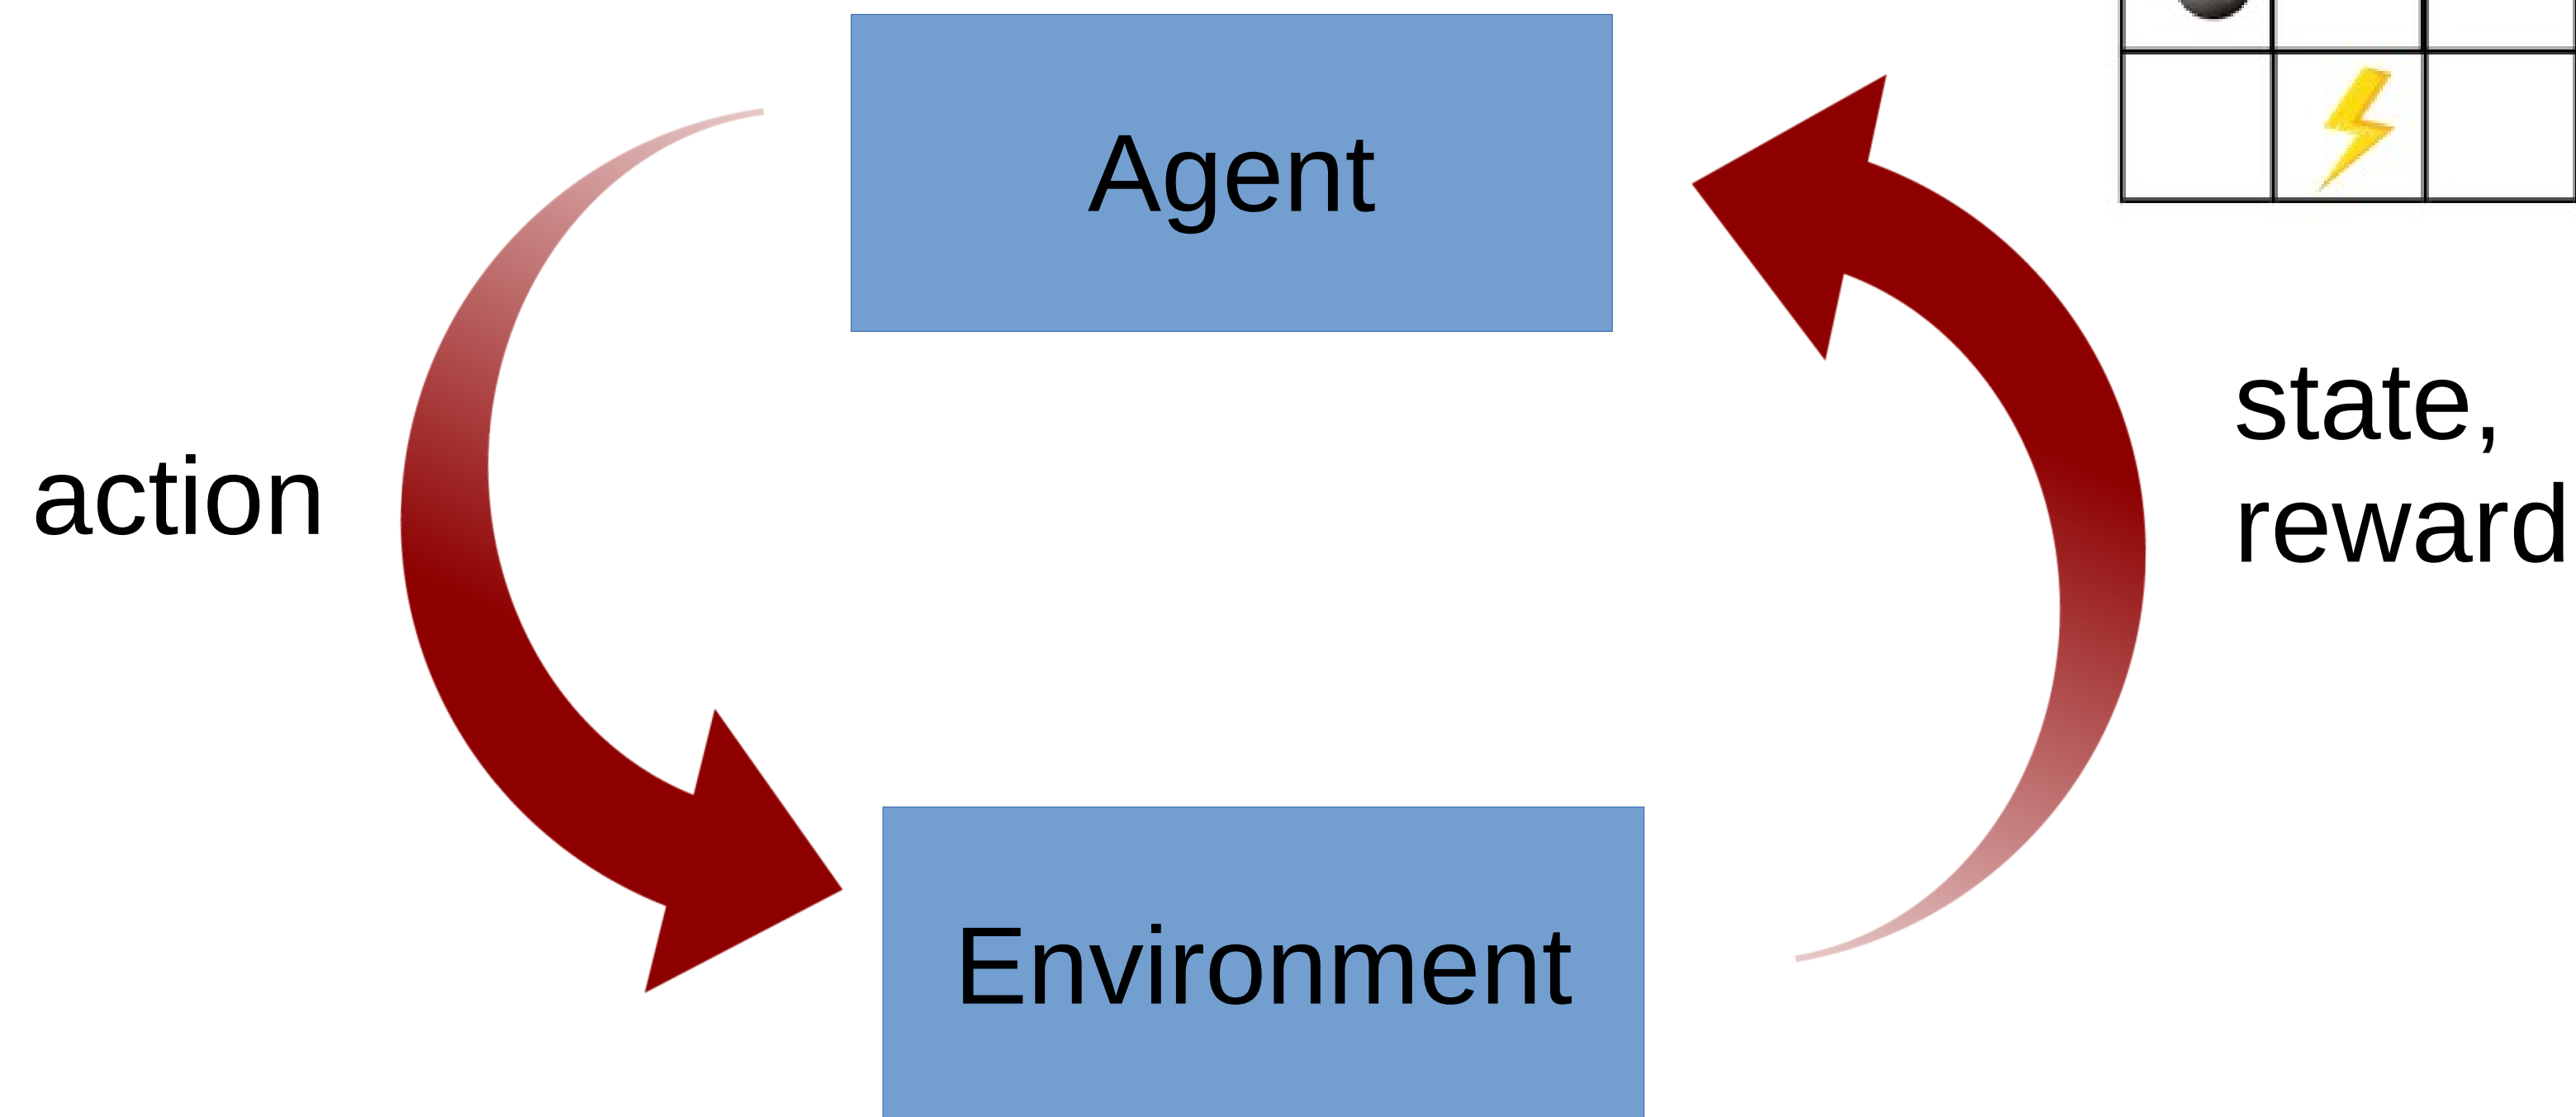
- Learning how to play games
- Robotics
- Finance
- Healthcare
- Meta-Learning



<https://ai.googleblog.com/2018/06/scalable-deep-reinforcement-learning.html>

Reinforcement Learning: Details


- Learning to maximize **rewards** by performing **actions** in an **environment**.



<https://www.freecodecamp.org/news/an-introduction-to-q-learning-reinforcement-learning-14ac0b4493cc/>

Reinforcement Learning: Details

- **Environment:** maze
- **Actions:** left/right/up/down
- **Rewards:**
 - -1 on each step,
 - -100 to step on mine
 - 1 for lightning charge
 - 100 for end
- **Policy:** mapping from states to actions, e.g.
 - always go left until wall, then right
 - after stepping on mine, always go right+down



<https://www.freecodecamp.org/news/an-introduction-to-q-learning-reinforcement-learning-14ac0b4493cc/>

Reinforcement Learning: Details

- **Q-Table:**

a table storing the expected rewards for every (state, action)-pair

Actions : ↑ → ↓ ←

Start				
Nothing / Blank				
Power				
Mines				
END				

<https://www.freecodecamp.org/news/an-introduction-to-q-learning-reinforcement-learning-14ac0b4493cc/>

Reinforcement Learning: Details

Scenario:

- a learning **agent**
- **S**: a set of possible states
- **A**: a set of possible actions
- a **state transition** function

$$\delta: S \times A \rightarrow S$$

- a **reward** function

$$r: S \times A \rightarrow \mathbb{R}$$

Feedback loop:

- the agent repeatedly chooses an action according to some **policy**

$$\pi: S \rightarrow A$$


- the environment changes to a new state according to δ
- some states provide the agent with feedback (**reinforcement**)

Reinforcement Learning: Reward

- Cumulative expected reward:

$$G_t = \sum_{i=0}^{\infty} \gamma^i * r_{t+i}$$


(γ makes the G_t finite)



- Cumulative expected reward:

$$G_t = \sum_{i=0}^{\infty} \gamma^i * r_{t+i}$$

(γ makes the G_t finite)



- Bellman **expectation** for the state-value function:


$$\begin{aligned} v_{\pi}(s) &= E[G_t | S_t = s] = E_{\pi}[R_t + \gamma * G_{t+1} | S_t = s] \\ &= \underbrace{\sum_a \pi(a|s)}_{\text{policy stochasticity}} \underbrace{\sum_{r, \hat{s}} p(r, \hat{s} | s, a)}_{\text{environment stochasticity}} * [r + \gamma * \underbrace{E_{\pi}[G_{t+1} | S_{t+1} = \hat{s}]}_{v_{\pi}(\hat{s})}] \end{aligned}$$

<https://www.coursera.org/learn/practical-rl/home/welcome>

<https://www.ke.tu-darmstadt.de/lehre/archiv/ss09/ki/reinforcement-learning.pdf>

- State-value to action-value function

$$v_{\pi}(s) = \sum_a \pi(a|s) q_{\pi}(s, a)$$



- Bellman **expectation** for the action-value function:

$$\begin{aligned} q_{\pi}(s, a) &= E[G_t | S_t = s, A_t = a] = E_{\pi}[R_t + \gamma * G_{t+1} | S_t = s, A_t = a] \\ &= \underbrace{\sum_{r, \hat{s}} p(r, \hat{s} | s, a)}_{\text{environment stochasticity}} * \underbrace{[r + \gamma * E_{\pi}[G_{t+1} | S_{t+1} = \hat{s}]]}_{v_{\pi}(\hat{s})} \end{aligned}$$

Reinforcement Learning: Optimal Policy

$$v_{opt}(s) = \max_{\pi} v_{\pi}(s)$$

$$\pi_{opt} = \arg \max_{\pi} v_{\pi}(s)$$

$$q_{opt}(s, a) = \max_{\pi} q_{\pi}(s, a)$$

$$\pi_{opt}(s) = \arg \max_a q_{\pi}(s, a)$$

Bellman **optimality** equations:

$$v_{opt}(s) = \max_a \underbrace{\sum_{r, \hat{s}} p(r, \hat{s} | s, a)}_{\text{environment stochasticity}} * [r + \gamma * v_{opt}(\hat{s})]$$

$$q_{opt}(s, a) = \underbrace{\sum_{r, \hat{s}} p(r, \hat{s} | s, a)}_{\text{environment stochasticity}} * [r + \gamma * \max_{\hat{a}} q_{opt}(\hat{s}, \hat{a})]$$

<https://www.coursera.org/learn/practical-rl/home/welcome>

Reinforcement Learning: Q-Learning

Model-free (train on trajectories),
Off-policy (not train on own policy)


$$\forall s \in S, \forall a \in A, q(s, a) = 0$$

Loop:

Sample $\langle s, a, r, \hat{s} \rangle$


Compute $\tilde{q}(s, a) = r(s, a) + \gamma * \max_{a_i} q(\hat{s}, a_i)$

Update $q(s, a) = \alpha * \tilde{q}(s, a) + (1 - \alpha) * q(s, a)$



<http://icaps18.icaps-conference.org/fileadmin/alg/conferences/icaps18/summerschool/lectures/Lecture5-rl-intro.pdf>
<https://www.coursera.org/learn/practical-rl/home/welcome>

Reinforcement Learning: Q-Learning




How to sample \hat{s} ?

ϵ -greedy policy

Exploration-exploitation trade-off:

With probability ϵ choose a **random** action,
else the **best** one.

Reinforcement Learning: Openai Gym



<https://gym.openai.com/envs/Taxi-v2/>

<https://www.coursera.org/learn/practical-rl/home/welcome>